

# Agenda

## Inspecting process and system state

More Tips Here: [https://www.cs.swarthmore.edu/~newhall/unixhelp/os\\_stats.php](https://www.cs.swarthmore.edu/~newhall/unixhelp/os_stats.php)

## File descriptors

## Signals

## Job Control

- Process Groups
- Sessions
- Terminal

# Inspecting process and system information

Linux stores kernel state in two pseudo-file systems

**/proc** – runtime info about processes (`man 5 proc` for info)

**/sys** – info about the system state (devices, buses, drivers, etc.)

/proc creates a directory for each process

- /proc/<pid>/fd // File descriptors
- /proc/<pid>/status // pid, session id, process group id
- /proc/<pid>/XXX

# /proc contains a directory for each process


```
alinen@sutekh:~/cs355/os-devel/lectures/processes$ ps -A
 PID TTY      TIME CMD
  1 hvc0    00:00:00 init(Ubuntu)
  5 hvc0    00:00:12 init
  8 ?
  9 ?
 10 pts/0   00:00:00 bash
 41 ?
 42 ?
 43 pts/1   00:00:00 bash
132 ?
133 ?
134 pts/2   00:00:00 bash
345 ?
346 ?
347 pts/3   00:00:00 bash
592 pts/0   00:00:00 ps
alinen@sutekh:~/cs355/os-devel/lectures/processes$
```

```
alinen@sutekh:/proc$ ls -l
total 0
dr-xr-xr-x  9 root  root
dr-xr-xr-x  9 alinen alinen
dr-xr-xr-x  9 root  root
dr-xr-xr-x  9 root  root
dr-xr-xr-x  9 alinen alinen
dr-xr-xr-x  9 root  root
dr-xr-xr-x  9 root  root
dr-xr-xr-x  9 alinen alinen
dr-xr-xr-x  9 root  root
dr-xr-xr-x  9 root  root
dr-xr-xr-x  9 alinen alinen
dr-xr-xr-x  9 root  root
dr-xr-xr-x  9 alinen alinen
dr-xr-xr-x  9 root  root
dr-xr-xr-x  9 root  root
dr-xr-xr-x  2 root  root
-r--r--r--  1 root  root
dr-xr-xr-x  4 root  root
0 Feb 10 08:35 1
0 Feb 10 08:35 10
0 Feb 10 09:04 132
0 Feb 10 09:04 133
0 Feb 10 08:38 134
0 Feb 10 09:11 345
0 Feb 10 09:11 346
0 Feb 10 09:11 347
0 Feb 10 09:04 41
0 Feb 10 09:04 42
0 Feb 10 08:35 43
0 Feb 10 08:36 5
0 Feb 10 11:06 598
0 Feb 10 09:04 8
0 Feb 10 09:04 9
0 Feb 10 11:03 acpi
0 Feb 10 11:03 buddyinfo
0 Feb 10 11:03 bus
```

directory  
is created  
for running  
process

# Inspecting process state

## Use **ps**



- ps -eo pid,ppid,tty,stat,cmd (Control what info is displayed)
- ps -efwH (Show process hierarchy)

## Use **top -u <username>**

Shows a configurable list of processes that updates dynamically

## Use **strace**

Displays all system calls made by an application

Displays any signals sent or received

# top

q to quit

f/F to filter/sort

v/V forest view

```
alinen@sutekh: ~/cs355/os-d × + ▾

top - 09:05:38 up 30 min, 0 users, load average: 0.00, 0.00, 0.00
Threads: 17 total, 1 running, 16 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.0 us, 0.1 sy, 0.0 ni, 99.9 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
MiB Mem : 15909.2 total, 15516.0 free, 325.2 used, 67.9 buff/cache
MiB Swap: 0.0 total, 0.0 free, 0.0 used. 15388.7 avail Mem

          PID USER      PR  NI    VIRT    RES    SHR   S %CPU %MEM TIME+ COMMAND
            1 root      20   0    2776   1956   1828 S  0.0  0.0  0:00.01 init(Ubuntu)
            7 root      20   0    2776   1956   1828 S  0.0  0.0  0:00.00 Interop
            5 root      20   0    3504   444    132 S  0.0  0.0  0:00.00 init
            6 root      20   0    3504   444    132 S  0.0  0.0  0:01.17 init
           274 root     20   0    3504   444    132 S  0.0  0.0  0:00.09 init
           287 root     20   0    3504   444    132 S  0.0  0.0  0:00.00 init
           288 root     20   0    3504   444    132 S  0.0  0.0  0:00.00 init
             8 root      20   0    2780   212     80 S  0.0  0.0  0:00.00 SessionLeader
             9 root      20   0    2780   216     80 S  0.0  0.0  0:00.00 Relay(10)
            10 alinen   20   0    6304   5280   3392 S  0.0  0.0  0:00.08 bash
            41 root      20   0    2780   212     80 S  0.0  0.0  0:00.00 SessionLeader
            42 root      20   0    2780   216     80 S  0.0  0.0  0:00.00 Relay(43)
            43 alinen   20   0    6344   5364   3476 S  0.0  0.0  0:00.02 bash
           132 root     20   0    2780   212     80 S  0.0  0.0  0:00.00 SessionLeader
           133 root     20   0    2780   216     80 S  0.0  0.0  0:00.00 Relay(134)
           134 alinen   20   0    6212   5076   3360 S  0.0  0.0  0:00.02 bash
           286 alinen   20   0    7824   3596   3008 R  0.0  0.0  0:00.03 top
```

# Demo: top H

```
alinen@sutekh: ~/cs355/os-d × + ▾ - □ ×

top - 12:22:54 up 14:46, 0 users, load average: 0.00, 0.00, 0.00
Threads: 13 total, 1 running, 12 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.0 us, 0.0 sy, 0.0 ni, 100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0
MiB Mem : 15909.2 total, 15356.9 free, 387.5 used, 164.7 buff/cach
MiB Swap: 0.0 total, 0.0 free, 0.0 used. 15277.8 avail Mem

PID USER S COMMAND PPID PGRP TTY TPGID
 6 root S `-- init          1   0 hvc0    0
 7 root S `-- Interop      0   0 hvc0    0
 8 root S `-- SessionLeader 1   8 ?     -1
 9 root S `-- Relay(10)     8   8 ?     -1
10 alinen S `-- bash       9  10 pts/0  5633
5633 alinen R `-- top      10  5633 pts/0  5633
741 root S `-- SessionLeader 1  741 ?     -1
742 root S `-- Relay(743)    741 741 ?     -1
743 alinen S `-- bash       742 743 pts/1  5640
5640 alinen S `-- fgbg      743 5640 pts/1  5640
5641 alinen S `-- fgbg      5640 5641 pts/1  5640
```

```
alinen@sutekh:~/cs355/os-devel/labs/03$ man sigprocm
alinen@sutekh:~/cs355/os-devel/labs/03$ vi fgbg.c
alinen@sutekh:~/cs355/os-devel/labs/03$ make
g++ -g -Wno-unused-variable -Wno-unused-but-set-vari
-o fgbg -lreadline
alinen@sutekh:~/cs355/os-devel/labs/03$ ./fgbg
[1] Created process 5641 in the background
$ |
```

# Demo: strace

Utility for tracking all system calls made by processes

To see all system calls made by a process and its children:

```
strace -f <exe>
```

To only see signals:

```
strace -e 'trace=!all' -f <exe>
```

To monitor dup2 of files: strace -e dup2 <exe>


To see all file related activity: strace -e trace=file <exe>

# Pipes and redirection


All processes have default file descriptors (fd) for stdin (0), stdout (1), stderr(2)

We can change the default input/output behavior by changing the file descriptors

ls | sort



ls > sort.txt



# Demo: Creating a pipe

```
int pipefd[2]; // 0 -> read; 1 -> write

int status = pipe(pipefd);
checkerror(status, "pipe");


pid_t cpid = fork();
checkerror(cpid, "fork");

if (cpid == 0) { // child
    close(pipefd[1]); // close unused write end
    while (read(pipefd[0], &buf, 1) > 0) {
        write(STDOUT_FILENO, &buf, 1);
    }
    write(STDOUT_FILENO, "\n", 1);
    close(pipefd[0]);
}
```


pipe is a System call  
initializes a pipe w/ a write end  
& a read end

```
else {
    close(pipefd[0]); // close unused read end
    const char* message = "mario takes the pipe to world 5";
    write(pipefd[1], message, strlen(message));
    close(pipefd[1]);
    wait(NULL); // wait for child
}
exit(EXIT_SUCCESS);
}
```

# Visualizing pipes



Parent writes to `fds[1]`  
Child reads from `fds[0]`



# Demo: pipes

```
alinen@sutekh:~/cs355/os-devel/lectures/processes$ ls -l /proc/1581/fd
total 0
lrwx----- 1 alinen alinen 64 Feb  4 20:15 0 -> /dev/pts/0
lrwx----- 1 alinen alinen 64 Feb  4 20:15 1 -> /dev/pts/0
lrwx----- 1 alinen alinen 64 Feb  4 20:15 2 -> /dev/pts/0
lr-x----- 1 alinen alinen 64 Feb  4 20:15 3 -> 'pipe:[5344]'
l-wx----- 1 alinen alinen 64 Feb  4 20:15 4 -> 'pipe:[5344]'

alinen@sutekh:~/cs355/os-devel/lectures/processes$ ls -l /proc/1582/fd
total 0
lrwx----- 1 alinen alinen 64 Feb  4 20:15 0 -> /dev/pts/0
lrwx----- 1 alinen alinen 64 Feb  4 20:15 1 -> /dev/pts/0
lrwx----- 1 alinen alinen 64 Feb  4 20:15 2 -> /dev/pts/0
lr-x----- 1 alinen alinen 64 Feb  4 20:15 3 -> 'pipe:[5344]'
l-wx----- 1 alinen alinen 64 Feb  4 20:15 4 -> 'pipe:[5344]'

alinen@sutekh:~/cs355/os-devel/lectures/processes$ ls -l /proc/1581/fd
total 0
lrwx----- 1 alinen alinen 64 Feb  4 20:15 0 -> /dev/pts/0
lrwx----- 1 alinen alinen 64 Feb  4 20:15 1 -> /dev/pts/0
lrwx----- 1 alinen alinen 64 Feb  4 20:15 2 -> /dev/pts/0
l-wx----- 1 alinen alinen 64 Feb  4 20:15 4 -> 'pipe:[5344]'

alinen@sutekh:~/cs355/os-devel/lectures/processes$ ls -l /proc/1582/fd
total 0
lrwx----- 1 alinen alinen 64 Feb  4 20:15 0 -> /dev/pts/0
lrwx----- 1 alinen alinen 64 Feb  4 20:15 1 -> /dev/pts/0
lrwx----- 1 alinen alinen 64 Feb  4 20:15 2 -> /dev/pts/0
lr-x----- 1 alinen alinen 64 Feb  4 20:15 3 -> 'pipe:[5344]'

alinen@sutekh:~/cs355/os-devel/lectures/processes$ |
```

# Demo: Redirecting output with dup2

```
void waitkey()
{
    printf("Press any key to continue"); getc(stdin);
    printf("\n");
}

int main(int argc, char *argv[])
{
    printf("Run `ls -l /proc/%d/fd` to see the file descriptors "
           "for this process.\n", getpid());
    waitkey();

    int fd = open("list.txt", O_RDONLY); waitkey();

    dup2(fd, STDERR_FILENO); waitkey();

    close(fd); waitkey();
    return 0;
}
```

\$ ./dup

Run `ls -l /proc/639/fd` to see the file descriptors for this process.  
Press any key to continue

# Visualizing dup2

```
int fd = open("list.txt", O_RDONLY);
dup2(fd, STDERR_FILENO);
close(fd);
```

# Demo: Dup2

```
alinen@sutekh:~$ ls -l /proc/639/fd
total 0
lrwx----- 1 alinen alinen 64 Feb 10 11:14 0 -> /dev/pts/0
lrwx----- 1 alinen alinen 64 Feb 10 11:14 1 -> /dev/pts/0
lrwx----- 1 alinen alinen 64 Feb 10 11:14 2 -> /dev/pts/0
alinen@sutekh:~$ ls -l /proc/639/fd
total 0
lrwx----- 1 alinen alinen 64 Feb 10 11:14 0 -> /dev/pts/0
lrwx----- 1 alinen alinen 64 Feb 10 11:14 1 -> /dev/pts/0
lrwx----- 1 alinen alinen 64 Feb 10 11:14 2 -> /dev/pts/0
lr-x----- 1 alinen alinen 64 Feb 10 11:14 3 -> /home/alinen/cs355/os-devel/lectures/processes/list.txt
alinen@sutekh:~$ ls -l /proc/639/fd
total 0
lrwx----- 1 alinen alinen 64 Feb 10 11:14 0 -> /dev/pts/0
lrwx----- 1 alinen alinen 64 Feb 10 11:14 1 -> /dev/pts/0
lr-x----- 1 alinen alinen 64 Feb 10 11:14 2 -> /home/alinen/cs355/os-devel/lectures/processes/list.txt
lr-x----- 1 alinen alinen 64 Feb 10 11:14 3 -> /home/alinen/cs355/os-devel/lectures/processes/list.txt
alinen@sutekh:~$ ls -l /proc/639/fd
total 0
lrwx----- 1 alinen alinen 64 Feb 10 11:14 0 -> /dev/pts/0
lrwx----- 1 alinen alinen 64 Feb 10 11:14 1 -> /dev/pts/0
lr-x----- 1 alinen alinen 64 Feb 10 11:14 2 -> /home/alinen/cs355/os-devel/lectures/processes/list.txt
alinen@sutekh:~$ |
```

NOTE: The special files in fd are **symbolic links**. A symbolic link refers to a file without copying it.

# Signals

**Signals** are a type of software interrupt. A small message to tell a process that some event has happened

How signals work:

1. OS sends a signal to a process

- On behalf of another process that called the `kill` syscall
- As the result of some event (NULL pointer dereference)

2. A process receives a signal

Asynchronous: signalee doesn't know when it will get one

Signals are pending before a process receives it

3. A signal interrupts the receiving process, which then runs signal handler code

- default handlers for each signal type in OS
- programmer can also add signal handler code

# Examples: Signal Types

<i>ID</i>	<i>Name</i>	<i>Default Action</i>	<i>Corresponding Event</i>
2	SIGINT	Terminate	Interrupt (e.g., <code>ctl-c</code> from keyboard)
9	SIGKILL	Terminate	Kill program (cannot override or ignore)
11	SIGSEGV	Terminate	Invalid memory reference (e.g. <code>NULL</code> ptr)
18	SIGCONT	Continue	Resume process (e.g. <code>bg/fg</code> shell commands)
19	SIGSTP	Stop	Stop process (e.g. <code>ctrl-z</code> from keyboard)
14	SIGALRM	Terminate	Timer signal
17	SIGCHLD	Ignore	Child stopped or terminated

segfault!

Signals can be sent explicitly (by a user) or implicitly (as a side-effect of program doing something, e.g. dereferencing a `NULL` ptr causes `SIGSEGV`)

``man 7 signal`` for more information

# Example: Cntrl-z, bg

```
$ ./inf_loop
^Z
[1]+  Stopped  ./inf_loop
$ ps w
  PID TTY      STAT      TIME  COMMAND
28850 pts/2    Ss       0:00  bash
29011 pts/2    T        0:03  ./inf_loop
29021 pts/2    R+       0:00  ps w
$ bg
./inf_loop
$ ps w
  PID TTY      STAT      TIME  COMMAND
28850 pts/2    Ss       0:00  bash
29011 pts/2    R        0:03  ./inf_loop
29105 pts/2    R+       0:00  ps w
```

Ctrl-z sends a SIGTSTP signal to every process running in the foreground, process is STOPPED

bg: sends SIGCONT signal and process runs in the background (shell continues)

ps w STAT field values:

*First letter:*

S: sleeping

T: stopped

R: running

*Second letter:*

s: session leader

+: foreground process

# Example: Cntrl-c, fg

```
$ ./inf_loop
^Z
[1]+  Stopped  ./inf_loop

$ ps w
  PID TTY      STAT      TIME  COMMAND
28850 pts/2    Ss       0:00  bash
29011 pts/2    T        0:03  ./inf_loop
29021 pts/2    R+       0:00  ps w

$ fg
./inf_loop
^C

$ ps w
  PID TTY      STAT      TIME  COMMAND
28850 pts/2    Ss       0:00  bash
29105 pts/2    R+       0:00  ps w
```

fg: sends SIGCONT  
and moves process  
into the foreground  
(shell waits)

Ctrl-c sends a  
SIGINT signal to  
every process  
running in the  
foreground  
(process will exit)

ps w STAT field values:

*First letter:*

S: sleeping

T: stopped

R: running

*Second letter:*

s: session leader

+: foreground process

# Sending signals

Signals can be sent using the **kill** command

Examples:

Shell command:

```
$ kill -9 1234      # send SIGKILL signal to process 1234  
$ kill -L            # List all signals
```

System call:

```
kill(1234, SIGKILL); // send SIGKILL to process 1234
```

# Example: Sending signals

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <signal.h>

int main() {

    int pid = 0;
    printf("Enter a process id: ");
    scanf(" %d", &pid);

    int signal_id = 0;
    printf("Enter a signal id: ");
    scanf(" %d", &signal_id);

    kill(pid, signal_id);
}
```

# Receiving signals

- A destination process *receives* a signal when it is forced by the kernel to react in some way to the delivery of the signal
- Possible ways to react:
  - **Ignore** the signal (do nothing)  
not all signals can be ignored (e.g. SIGKILL, SIGSTP)
  - **Terminate** the process on receipt of signal, possibly with **core** file
  - **Catch** the signal by executing a user-level function called signal handler

# Signal Handlers

To catch a signal, a program needs to register a signal handler

Modifies the default action associated with the receipt of a particular signal

A process can register a function (e.g. the signal handler) to be invoked when a signal is received

**#include <sys signal.h>** lists signal handling functions

To define a fn ptr : ~~type~~ typedef void (\*handler\_t)(int);

## Writing your own signal handler

```
signal(int signum, handler_t *handler);
```

↑ function ptr's type is determined by its return value & parameters

- handler is a function, indicated with a function pointer
  - When program receives signal, it jumps to start executing the handler function.
  - When the handler done executing, control passes back to instruction in the control flow of the process that was interrupted by receipt of the signal

# Example: Simple signal handler

```
// signal handler function: called when process receives SIGINT
#include <stdio.h>
#include <unistd.h>
#include <signal.h>
#include <stdlib.h>
#include <sys/wait.h>

void int_handler(int sig) {
    printf("Proc %d received signal %d\n",getpid(), sig);
    exit(0);
}

void main() {
    signal(SIGINT, int_handler);
    while(1);
}
```

# Job Control

Chapter 34 from TLPI

*The Linux Programming Interface*, by  
Michael Kerrisk, No Starch Press, 2010

Lots of examples

<https://www.man7.org/tlpi/>

## THE LINUX PROGRAMMING INTERFACE

A Linux and UNIX\* System Programming Handbook

MICHAEL KERRISK



# Job Control Overview

Modern shells support **job control** features that allow the user to start, pause and stop processes running at the command line

First introduced in 1980 on BSD (Kerrisk 2010)

A **job** is a command, typed at the command line. It can contain pipes, file redirections, and multiple commands, e.g.

```
ls -l | wc -w  
./compute_stats > output.txt  
./generate_data | ./compute_stats > output.txt
```

A job can run in the **foreground** (receives terminal input) or **background** (Demo)

# Terminal

The **terminal** corresponds to the window where we type in commands and see output

When you run a process in the foreground, the file descriptors – `stdout` (1), `stdin` (0), `stderr` (2) – are configured to write to the terminal

Writing to the terminal involves writing to a virtual file (**/dev/tty**) that represents the interface to our keyboard/mouse/monitor

# Example: Job Control at the terminal

```
$ echo $$                                Display the PID of the shell
400
$ find / 2> /dev/null | wc -l &        Creates 2 processes in background group
[1] 659
$ sort < longlist | uniq -c          Creates 2 processes in foreground group
```

# tty

tty = Teletypewriter

Each terminal gets a tty for reading and writing

```
$ ps w
```

PID	TTY	STAT	TIME	COMMAND
10	pts/0	Ss	0:00	-bash
43	pts/1	Ss+	0:00	-bash
134	pts/2	Ss+	0:00	-bash
347	pts/3	Ss+	0:00	-bash
738	pts/0	R+	0:00	ps w



# Demo: tty

Writing to the console corresponds to writing to the console's TTY

The screenshot shows a terminal window with four identical command entries:

```
alinen@sutekh:~$ echo "TEST" > /dev/pts/0
alinen@sutekh:~$ echo "TEST" > /dev/pts/0
alinen@sutekh:~$ echo "TEST" > /dev/pts/0
alinen@sutekh:~$
```

Below this, another terminal window is open in the same session, showing the output of the command:

```
alinen@sutekh:~/cs355/os-devel/lectures/processes$ TEST
TEST
TEST
TEST
TEST
TEST
TEST
TEST
TEST
```

A callout box contains the following warning:

WARNING: Processes should only read/write to its **controlling terminal**. Weird things can happen when you start messing with the virtual terminal files!!

# Who owns the terminal?

Foreground process groups (i.e., Foreground Jobs)

- Can read from STDIN

- Receive signals from the keyboard (e.g., CTRL + C)

To move a process from the background to the foreground, you  
“give” the background process the terminal

```
int tcsetpgrp(int fd, pid_t pgrp) // fd = STDIN_FILENO
```

# Watch out!

Only the foreground process can give the terminal to another process group

If a background process tries to take control of the terminal with **tcsetpggrp**, then the group gets sent **SIGTTOU**, which will stop the process group

If a background process tries to read from stdin, it gets sent the signal **SIGTTIN**

In your shell, the parent process should call **tcsetpggrp** to move processes from the foreground and background => the shell should ignore the signals SIGTTIN and SIGTTOU(More later)

Aside: Writing to stdout from the background is ok, but can be configured so that background processes get SIGTTOU

# Sessions and process groups

Each job belongs to the same **process group**

By default, the process group ID (**PGID**) corresponds to the **process group leader**. The process group leader is typically the parent when using fork and exec.

A **session** is a collection of process groups

The session id (**SID**) corresponds to the process ID of the session leader.

Associated with a login or a terminal

The session leader is typically the shell program that is run when the terminal starts

All processes in a session have a single **controlling terminal**

Within a session, a single process group is the **foreground process group**. Any number of process groups can be a **background process group**

# Example: Session and process groups

```
$ echo $$  
400  
$ find / 2> /dev/null | wc -l &  
[1] 659  
$ sort < longlist | uniq -c
```




Figure 34-1: Relationships between process groups, sessions, and the controlling terminal

# Example: Session and process groups

```
$ echo $$  
400  
$ find / 2> /dev/null | wc -l &  
[1] 659  
$ sort < longlist | uniq -c
```




Figure 34-1: Relationships between process groups, sessions, and the controlling terminal

# Example: Job control in the shell

```
$ grep -r SIGHUP /usr/src/linux >x &  
[1] 18932  
$ sleep 60 &  
[2] 18934
```

*Job 1: process running grep has PID 18932*

*Job 2: process running sleep has PID 18934*

```
$ jobs  
[1]- Running      grep -r SIGHUP /usr/src/linux >x &  
[2]+ Running      sleep 60 &
```

```
$ fg %1  
grep -r SIGHUP /usr/src/linux >x
```

*Job 1: move to foreground*

```
[1]- Done        grep -r SIGHUP /usr/src/linux >x  
[2]+ Done        sleep 60
```

*Completed jobs are reported when finished*

# Process groups

When a signal is sent, it sent to the entire process group by default

When we create a process with fork(), the child belongs to the same process group as the parent

The process group ID is equal to a process ID

- The process ID of the first process to exist in the group

Question: If a process group “leader” terminates, can its process ID be reused by another process? Even if the old group is still going?

- Answer: no, that process ID will be reserved until the group is done

# Setting the process group

```
int setpgid(pid_t pid, pid_t pgid)
```

- Only works if pgid specifies an existing process group
- Or if pgid == pid, creates a new process group of that id

```
pid_t getpgid(pid_t pid)
```

- Gets the process group id of the specified process
- If 0 is passed in, get the group ID of the calling process

# Process group example


```
// Putting the child in its own process group
```

```
int pid = fork();
if (pid == 0) {
    setpgid(0, 0); // put child in their own group; can also do setpgid(getpid(), getpid())
    while(1){ pause(); }
}
else {
    setpgid(pid, pid); // put child in their own group

    std::string command;
    std::cin >> command;
    while (command != "quit") std::cin >> command;

    kill(pid, SIGTERM); // make sure to kill child
}
```


# Example: Ctrl-C Same Group



SIGINT is sent to  
every  
Process in the  
process  
group

# Example: Ctrl-C different groups

**CTRL +C, different group**



SIGINT is sent to every Process in the process group

Child is in a separate group

# Waiting on processes

Can pass in -PGID (negative PGID) to kill() and waitpid()

Doing so for kill() will send the signal to all processes in the group

Doing so for waitpid() will wait for any process in the group